

fortnine

July 31, 2024

Author: Sraddha Bhattacharjee

Mount drive path

```
[29]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Question 1 solution

```
[30]: import pandas as pd

# Load the sales data
df = pd.read_csv("/content/drive/MyDrive/sales.csv", parse_dates=['orderdate'])

# Step 1: Identify the first order date for each customer
customer_first_order = df.groupby('customer_number')['orderdate'].min().
    ↪reset_index()
customer_first_order.columns = ['customer_number', 'first_order_date']

# Step 2: Merge the first order date back into the original dataframe
df = pd.merge(df, customer_first_order, on='customer_number')

# Step 3: Classify each order as new or returning
df['order_year'] = df['orderdate'].dt.year
df['customer_type'] = df.apply(lambda row: 'New' if row['orderdate'].year ==
    ↪row['first_order_date'].year else 'Returning', axis=1)

# Step 4: Count the number of new and returning customers per year
yearly_customer_count = df.groupby(['order_year',
    ↪'customer_type'])['customer_number'].nunique().reset_index()
yearly_customer_count.columns = ['Year', 'CustomerType', 'CustomerCount']

# Step 5: Calculate the total number of customers per year
total_customers_per_year = yearly_customer_count.
    ↪groupby('Year')['CustomerCount'].sum().reset_index()
total_customers_per_year.columns = ['Year', 'TotalCustomers']
```

```

# Step 6: Merge the total customers per year back into the yearly customer_
↳count dataframe
yearly_customer_count = pd.merge(yearly_customer_count,
↳total_customers_per_year, on='Year')

# Step 7: Calculate the portion of new and returning customers
yearly_customer_count['Portion'] = yearly_customer_count['CustomerCount'] /
↳yearly_customer_count['TotalCustomers']

# Step 8: Pivot the dataframe to get the desired format
result = yearly_customer_count.pivot(index='Year', columns='CustomerType',
↳values='Portion').fillna(0)

# Step 9 : Add new + returning column values which always add to 1

result['Total'] = result.sum(axis=1)

print(result)

```

CustomerType	New	Returning	Total
Year			
2018	1.000000	0.000000	1.0
2019	0.660721	0.339279	1.0
2020	0.624806	0.375194	1.0
2021	0.468724	0.531276	1.0
2022	1.000000	0.000000	1.0
2023	0.516129	0.483871	1.0

Question 2 solution:

```

[31]: import pandas as pd

# Step 1: Load the data from the CSV file
df = pd.read_csv("/content/drive/MyDrive/sales.csv")

# Step 2: Identify and Remove Outliers using the IQR method
def remove_outliers(group):
    Q1 = group['QuantityOrdered'].quantile(0.25)
    Q3 = group['QuantityOrdered'].quantile(0.75)
    IQR = Q3 - Q1
    return group[(group['QuantityOrdered'] >= (Q1 - 1.5 * IQR)) &
↳(group['QuantityOrdered'] <= (Q3 + 1.5 * IQR))]

df_no_outliers = df.groupby('sku_id').apply(remove_outliers).
↳reset_index(drop=True)

```

```
[41]: import pandas as pd

# Convert 'orderdate' column to datetime format
df_no_outliers['orderdate'] = pd.to_datetime(df_no_outliers['orderdate'])

# Extract the year from 'orderdate' and create a new column 'order_year'
df_no_outliers['order_year'] = df_no_outliers['orderdate'].dt.year

# Group by 'order_year' and 'sku_id' to calculate average demand
# For each SKU and year, calculate the average quantity ordered
average_demand = df_no_outliers.groupby(['order_year', 'sku_id'])['QuantityOrdered'].mean().reset_index()

# Display the DataFrame with average demand per year for each SKU
print(average_demand)
```

	order_year	sku_id	QuantityOrdered
0	2018	1109285499	1.375000
1	2018	1109285500	4.000000
2	2018	1109285501	3.700000
3	2018	1109285502	1.000000
4	2018	1109285503	1.312883
...
73094	2023	1109317660	1.000000
73095	2023	1109317783	1.000000
73096	2023	1109381246	1.000000
73097	2023	1109382949	1.000000
73098	2023	1109382953	1.000000

[73099 rows x 3 columns]

```
[42]: # Shows number of unique rows in the output
len(average_demand.sku_id.unique())
```

[42]: 22743

Running Linear Regression model to predict next year's demand

```
[35]: import pandas as pd
from sklearn.linear_model import LinearRegression
import numpy as np

# Linear regression function to predict next year's demand
def predict_next_year_demand(df):
    results = []
    for sku_id, group in df.groupby('sku_id'):
        X = group['order_year'].values.reshape(-1, 1)
```

```

y = group['QuantityOrdered'].values

if len(group) < 2: # Not enough data points for a linear model
    predicted_demand = y[0] # Use the same quantity for the forecast
else:
    # Train the model
    model = LinearRegression()
    model.fit(X, y)

    # Predict the demand for the next year (max year + 1)
    next_year = group['order_year'].max() + 1
    predicted_demand = model.predict(np.array([[next_year]]))[0]

results.append({'sku_id': sku_id, 'quantity': predicted_demand})

return pd.DataFrame(results)

# Predict the next year's demand for each sku_id
predicted_demand_df = predict_next_year_demand(average_demand)
predicted_demand_df

```

```

[35]:
      sku_id  quantity
0      81588  1.000000
1    710294272  1.000000
2    1109285499  1.180147
3    1109285500  0.425000
4    1109285501  7.644505
...
22738  1109410606  1.000000
22739  1109410610  1.000000
22740  1109410613  1.000000
22741  1109411156  3.000000
22742  1109411157  1.000000

```

[22743 rows x 2 columns]

As an additional step I created predicted_demand_adjusted table which adjusts the predicted quantities:

- a) Sets values between 0 and 1 to 1.
- b) Sets non-positive values to 0.
- c) Leaves other values unchanged.

To ensure practical forecasted stock levels, I adjusted the Linear Regression model's predictions by setting quantities to zero or one when they are close to zero, preventing unrealistic negative forecasts.

```
[39]: # Adjust the quantities
predicted_demand_adjusted = predicted_demand_df.copy()
predicted_demand_adjusted['quantity'] = predicted_demand_adjusted['quantity'].
    ↪ apply(
        lambda x: 1 if 0 < x < 1 else (0 if x <= 0 else x)
    )

# Display the adjusted DataFrame
predicted_demand_adjusted
```

```
[39]:      sku_id  quantity
0      81588  1.000000
1    710294272  1.000000
2    1109285499  1.180147
3    1109285500  1.000000
4    1109285501  7.644505
...
22738  1109410606  1.000000
22739  1109410610  1.000000
22740  1109410613  1.000000
22741  1109411156  3.000000
22742  1109411157  1.000000
```

[22743 rows x 2 columns]

Reading the final data with sku_id and adjusted quantities to a csv named 'pred_demand'.

```
[40]: output_file_path = '/content/drive/MyDrive/pred_demand.csv' # Update this to
    ↪ the desired output path
predicted_demand_adjusted.to_csv(output_file_path, index=False)
```

Plotting the average demand graph for a random sku_id (= 1109286010) with data from previous years along with data predicted by the Linear Regression model for the next year

```
[60]: sample_row = average_demand[average_demand['sku_id'] == 1109286010]
sample_row
```

```
[60]:      order_year      sku_id  QuantityOrdered
224      2018  1109286010      145.777778
16221     2019  1109286010      159.933333
33799     2020  1109286010      194.484848
53901     2021  1109286010      220.520000
```

```
[64]: sample_row_predicted_value =
    ↪ predicted_demand_adjusted[predicted_demand_adjusted['sku_id'] == 1109286010]
```

```
sample_row_predicted_value['order_year'] = 2022
sample_row_predicted_value.columns = ['sku_id', 'QuantityOrdered', 'order_year']
sample_row_predicted_value
```

<ipython-input-64-01ea5566cbdd>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
sample_row_predicted_value['order_year'] = 2022

```
[64]:      sku_id  QuantityOrdered  order_year
      241  1109286010          244.873535      2022
```

```
[65]: plot_df = pd.concat([sample_row, sample_row_predicted_value])
      plot_df
```

```
[65]:      order_year      sku_id  QuantityOrdered
      224          2018  1109286010          145.777778
      16221         2019  1109286010          159.933333
      33799         2020  1109286010          194.484848
      53901         2021  1109286010          220.520000
      241          2022  1109286010          244.873535
```

```
[ ]:
```

```
[70]: import matplotlib.pyplot as plt
      # Create the scatter plot
      plt.figure(figsize=(10, 6))

      # Plot all points
      plt.scatter(plot_df['order_year'], plot_df['QuantityOrdered'], color='blue',
                  ↪label='Actual values')

      # Highlight the year 2022 with a different style
      df_2022 = plot_df[plot_df['order_year'] == 2022]
      plt.scatter(df_2022['order_year'], df_2022['QuantityOrdered'], color='red',
                  ↪marker='x', s=100, label='Predicted value')

      # Fit a linear regression model
      X = plot_df[['order_year']].values
      y = plot_df['QuantityOrdered'].values
      model = LinearRegression()
      model.fit(X, y)

      # Generate points for the regression line
```

```

x_range = np.linspace(X.min(), X.max(), 100).reshape(-1, 1)
y_pred = model.predict(x_range)

# Plot the regression line
plt.plot(x_range, y_pred, color='green', linestyle='--', label='Regression_
↳line')

# Add labels and title
plt.xlabel('Year')
plt.ylabel('Quantity Ordered')
plt.title('Scatter Plot of Quantity Ordered vs Year for sku_id = 1109286010 ')
plt.legend()

# Set x-ticks to only show whole numbers
plt.xticks(np.arange(plot_df['order_year'].min(), plot_df['order_year'].max() +
↳1, 1))

# Show the plot
plt.grid(True)
plt.show()

```

